

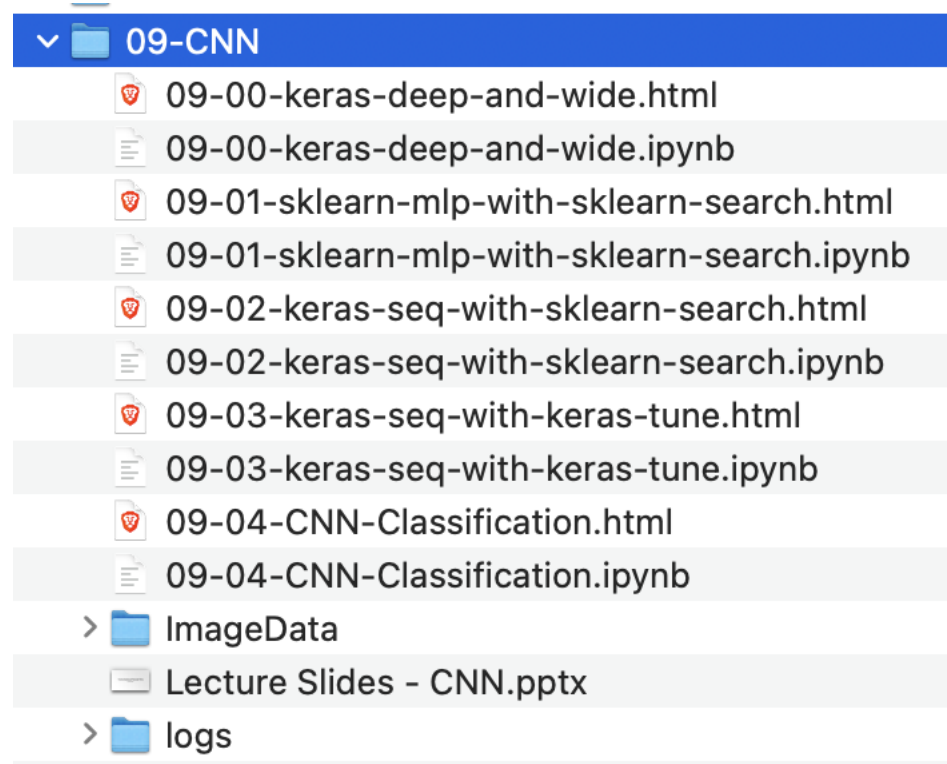
Convolutional Neural Nets

Tim Smith, PhD

Class09

- Review DNN
 - Demonstrate various deep and wide modeling techniques
- Discuss and elaborate on hyperparameter tuning
 - Sklearn model with Sklearn tuning
 - Keras model with Sklearn tuning
 - Keras model with Keras tuning
- Introduce the concept of convolutional neural networks
 - Build an image recognition model using CNN
- Demonstrate debugging models using Tensorflow

For the first part of the lecture:



Download the files from the following link...

<https://usf.box.com/s/7bn9cox3sqyyw32lqqps2cb4nmcq0buw>

Convolutional Neural Networks

- "Discrete convolution": *a mathematical operation on two functions (f and g) to produce a third function that expresses how the shape of one is modified by the other (Wikipedia)*
- Input data: **images (usually)**
- Used in
 - Image search
 - Image classification
 - Self-driving cars

Review on Images

- If black and white, there is 1 channel
- If color, there are 3 channels: Red, Green, Blue (RGB)





Image Data

- Image size is measured in pixels.
- Ex: 480 x 480 means:






Image Data

- Each pixel has a value between 0-255



28 x 28
784 pixels

[illegible]

Image Data

- Let's think of a 1M (megapixel) image
- 1 megapixel = 1,000,000 pixels (1,000 x 1,000)


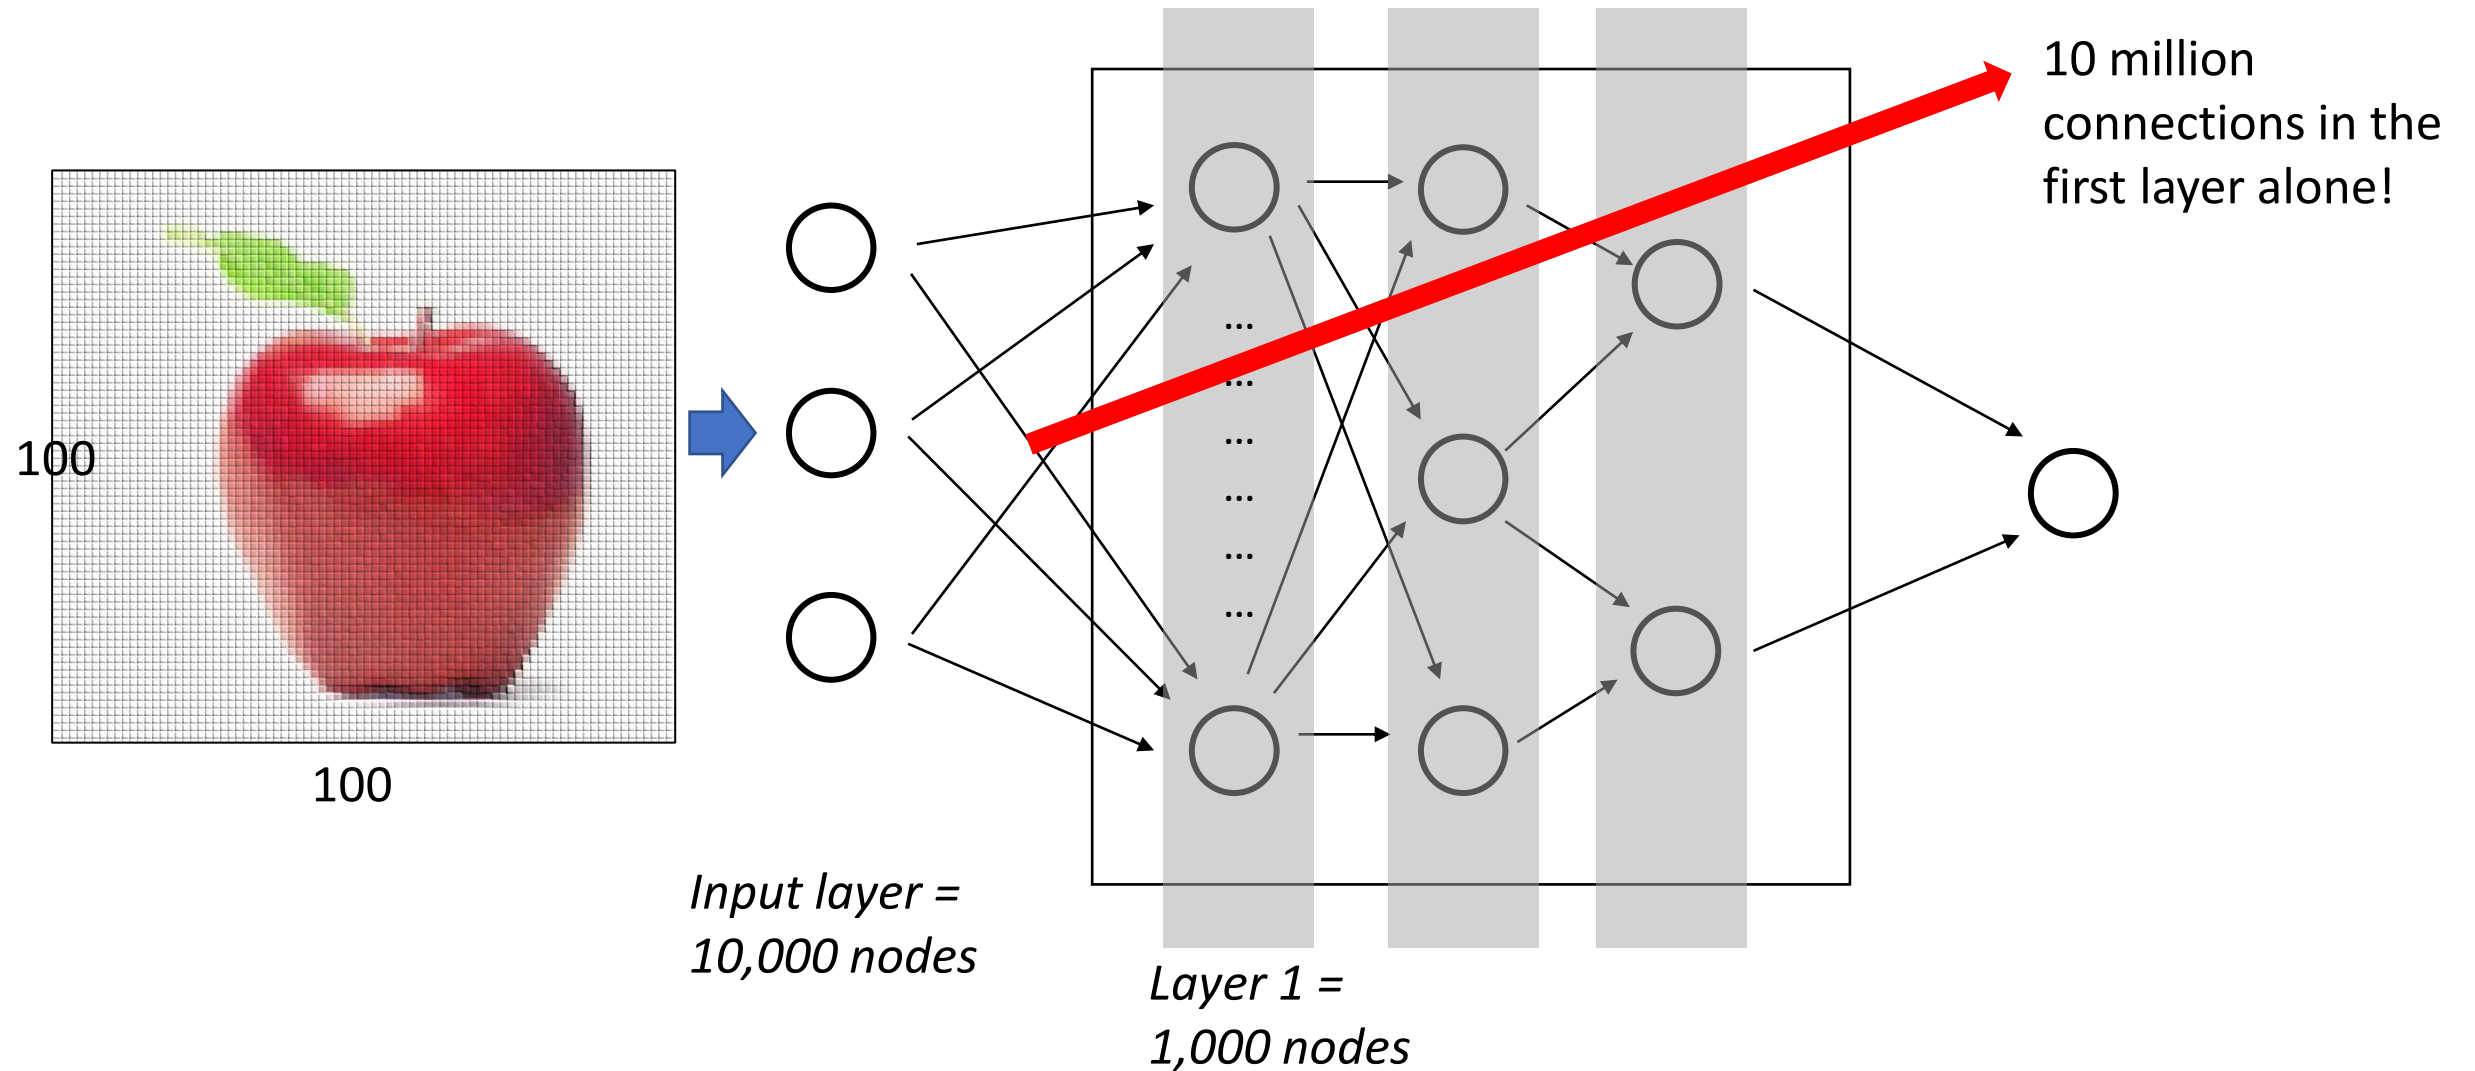


Image Data




Solution:

- **STEP 1 (convolution):** Extract features from an image:
 - Horizontal lines,
 - Vertical lines,
 - Curvature,
 - Colors,
 - Etc.
- **STEP 2 (pooling):** Reduce the variables (without losing information)
- **GET HELP FROM HUMAN VISUAL CORTEX**


Convolutional Neural Networks

- The visual cortex is made up of neurons
- Each neuron has a specific receptive field (and do not necessarily see the things outside of that field)



Convolutional Neural Networks

- Further, some neurons react only to horizontal lines
- Some react to different orientations
- Higher level neurons make sense of the pattern detected at lower levels




Convolutional Layer

- Apply the same logic:

“Receptive field”

Also called:

- “**Filter**”
- “Kernel”




Convolutional Layer

- Apply the same logic:

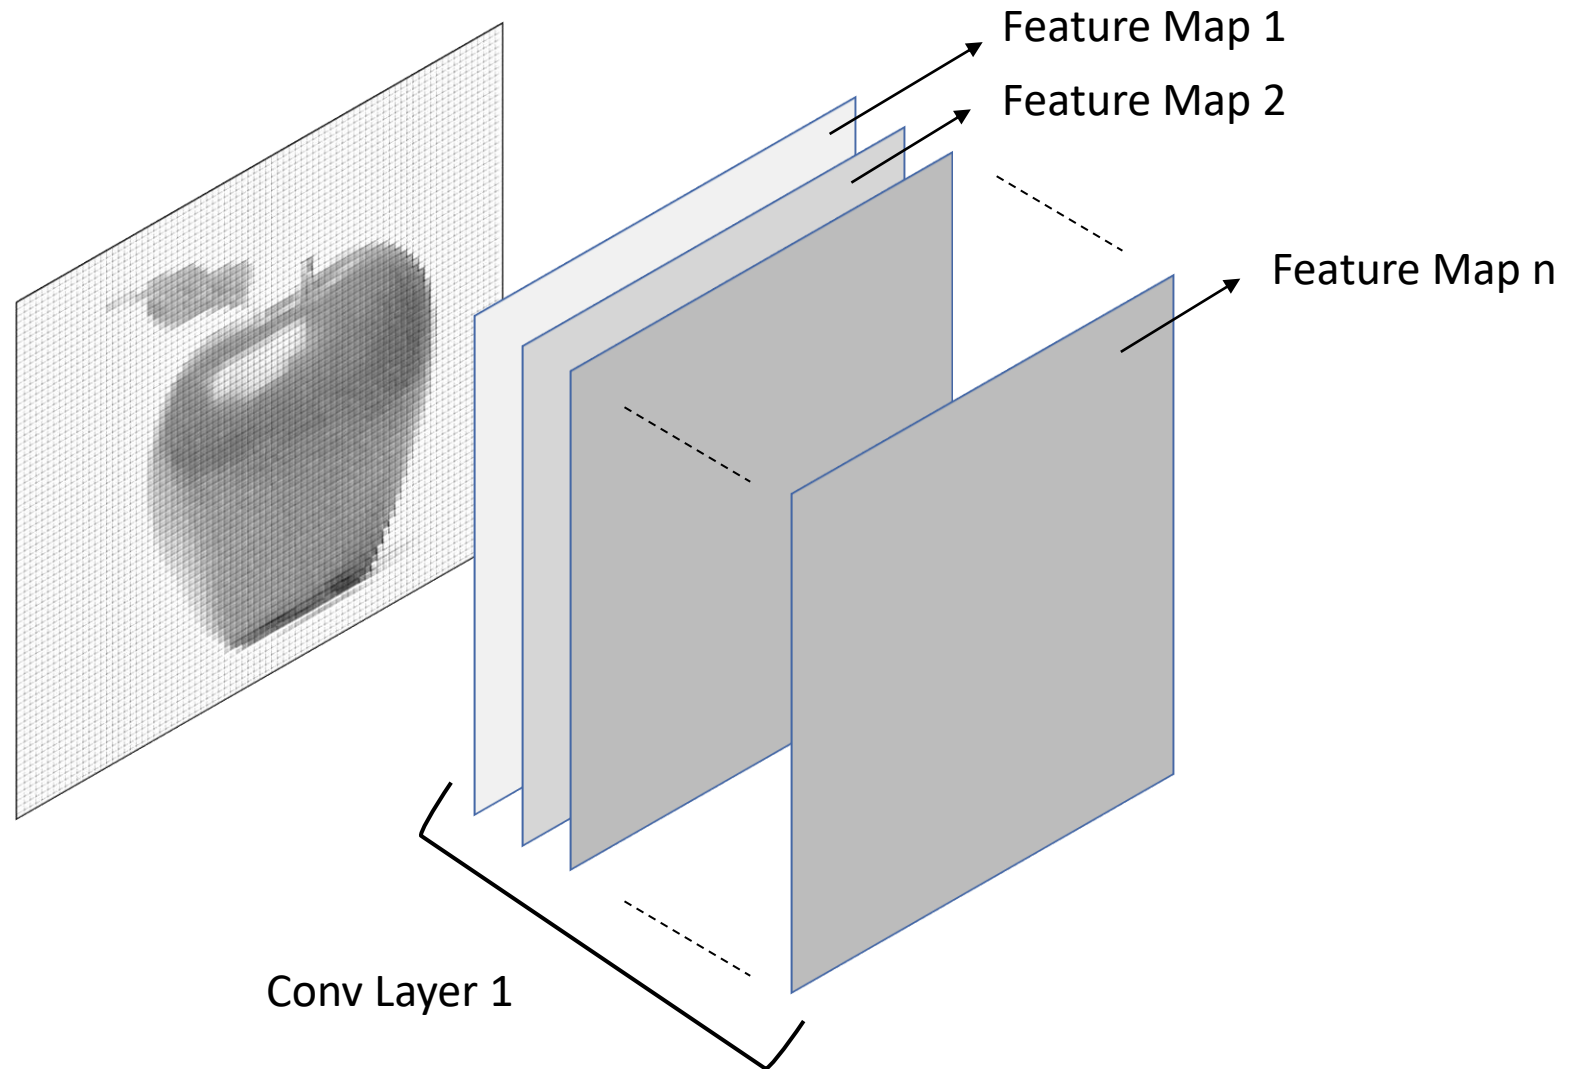
“Receptive field”

Also called:


- “**Filter**”
- “Kernel”



Convolutional Layer




Convolutional Layer



Makes the size of second layer the same as first


Convolutional Layer

- You can have varying strides (both horizontally, and vertically)




Convolutional Layer

- If you have RGB (3-channel) image, the filter has 3 channels too



Convolutional Layer

- Eventually, you have multiple feature maps (per convolutional layer)



Convolutional Layer


- How to process the values inside the “filter”
- Multiply them with “weights” then add them up
- Example: Apply “vertical” filter (to highlight vertical lines)

5 x 5 image

0	1	0	1	0
0	1	0	1	0
0	1	1	1	0
0	0	0	1	0
0	0	0	1	0

3 x 3 filter

0	1	0
0	1	0
0	1	0



Convolutional Layer

- Then, apply “horizontal” filter (to highlight horizontal lines)

5 x 5 image

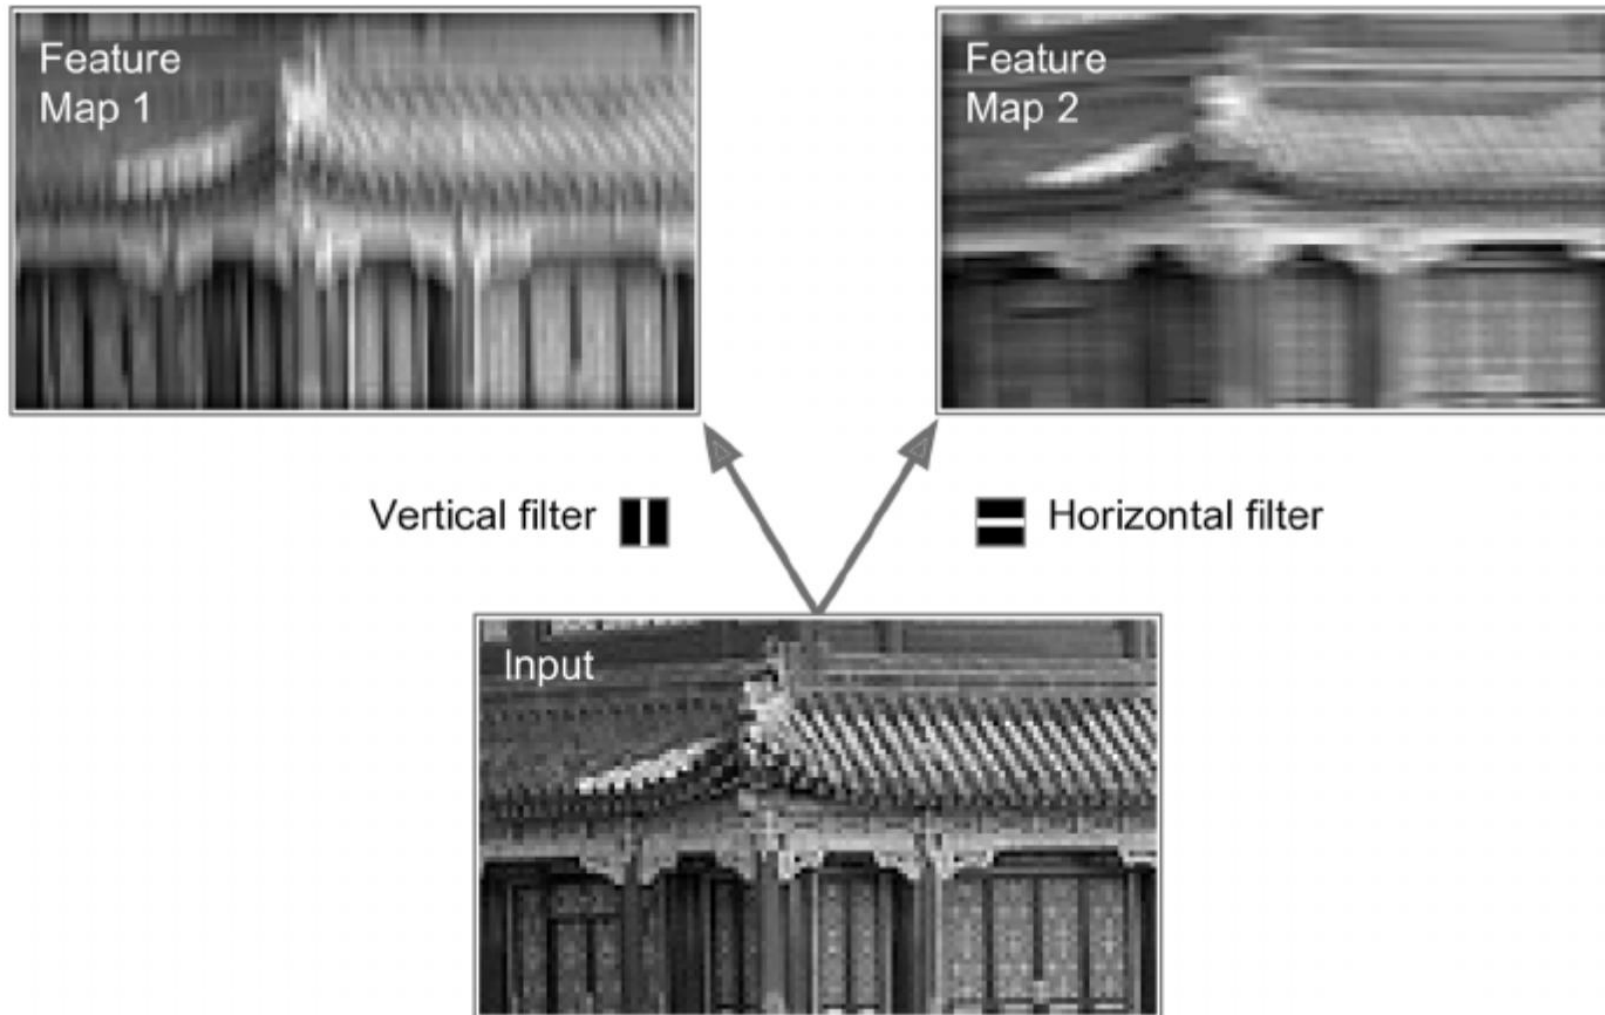
0	1	0	1	0
0	1	0	1	0
0	1	1	1	0
0	0	0	1	0
0	0	0	1	0

3 x 3 filter

0	0	0
1	1	1
0	0	0

- There are many more filters...

Convolutional Layer



Memory Requirements


- Imagine a single image with 150 x 100 pixels
 - Filter of 5 x 5 with 1 stride
 - 200 feature maps
 - Padding (add when necessary)
- Calculations:
 - $(5 \times 5 \times 3 + 1) \times 200 = 15,200$ terms
 - Each feature map has 150 x 100 neurons that need to compute 5 x 5 x 3 inputs. Makes 225 million multiplications
- 1 image needs 12 MB RAM.
- 100 images 1.2 GB RAM.

Pooling Layer

- “Shrinks” the convolutional layer (by sampling)
- An “aggregation” technique!
- Why:
 - Reduce computational load
 - Reduce number of parameters
 - Reduce overfitting

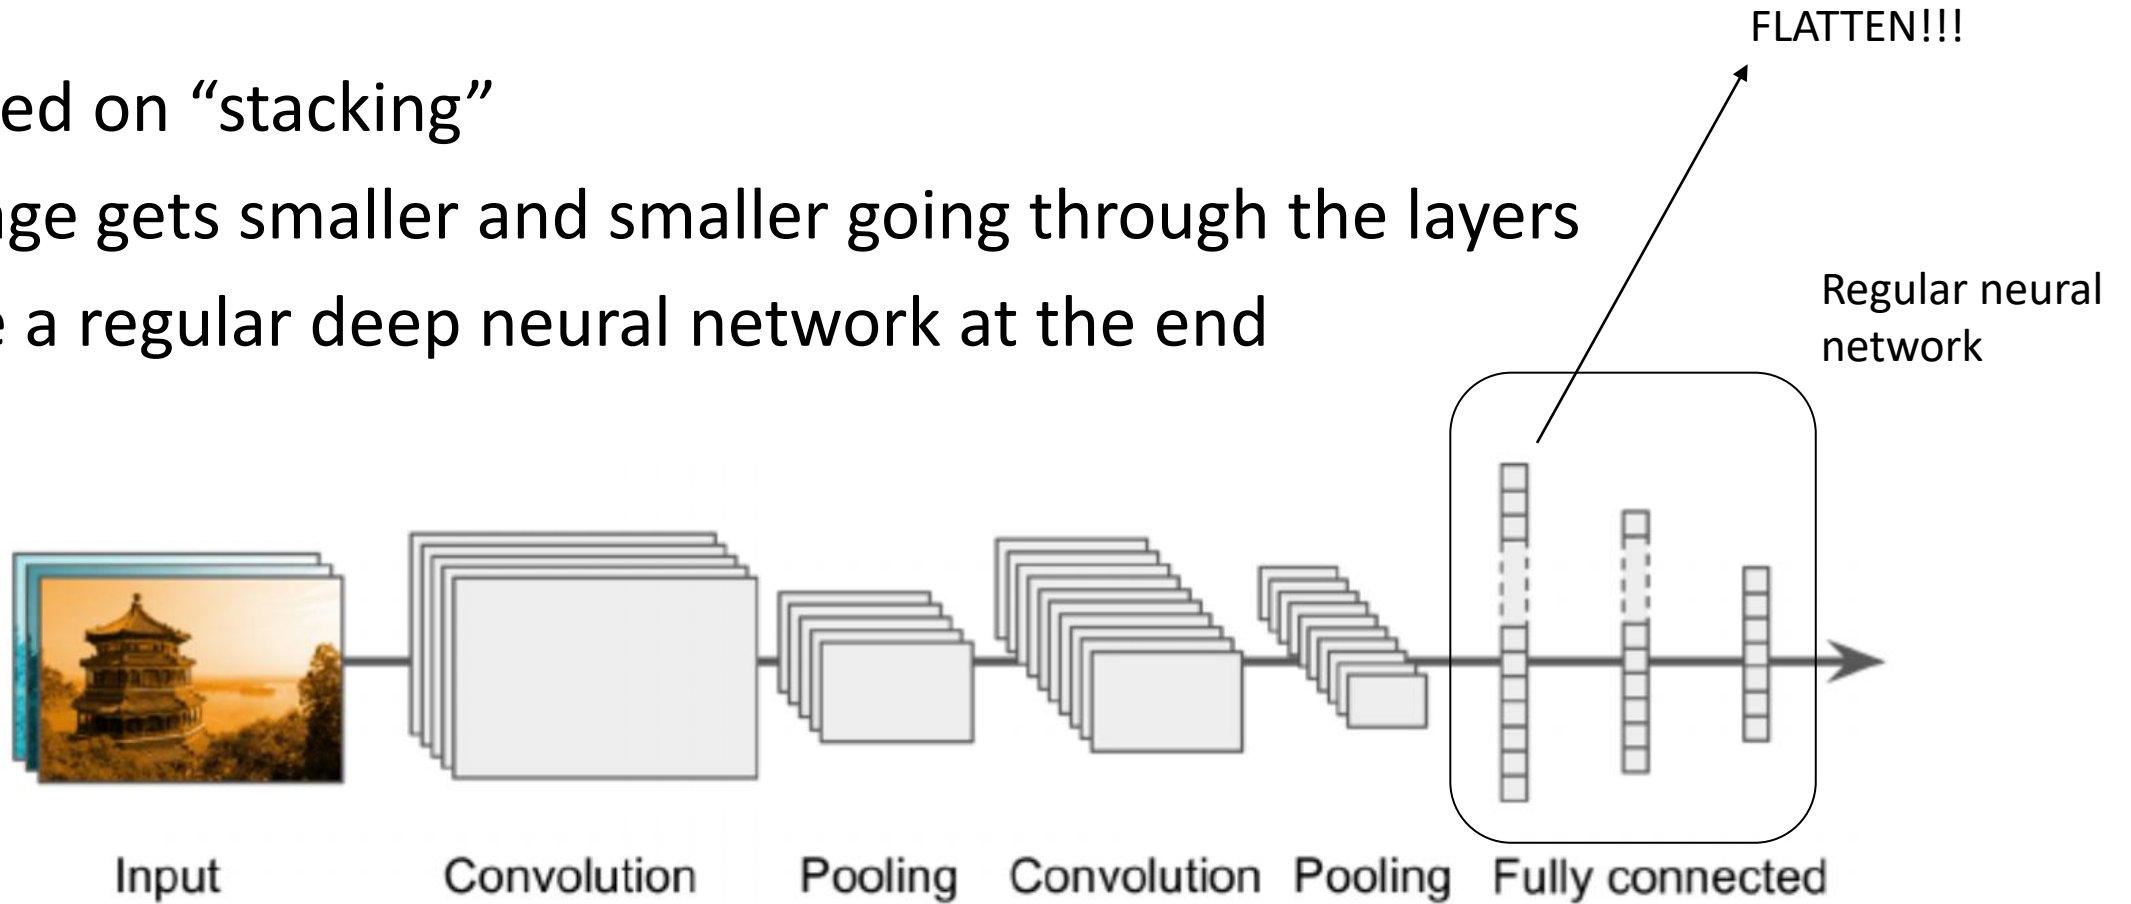
Pooling Layer

- Similar to earlier: define the size, stride, padding type
- It AGGREGATES the values
 - Example: use min, max, average, etc.
- A 2 x 2 kernel with a stride of 2 reduces the image (inputs) by 75%
- You can set it to shrink channels too



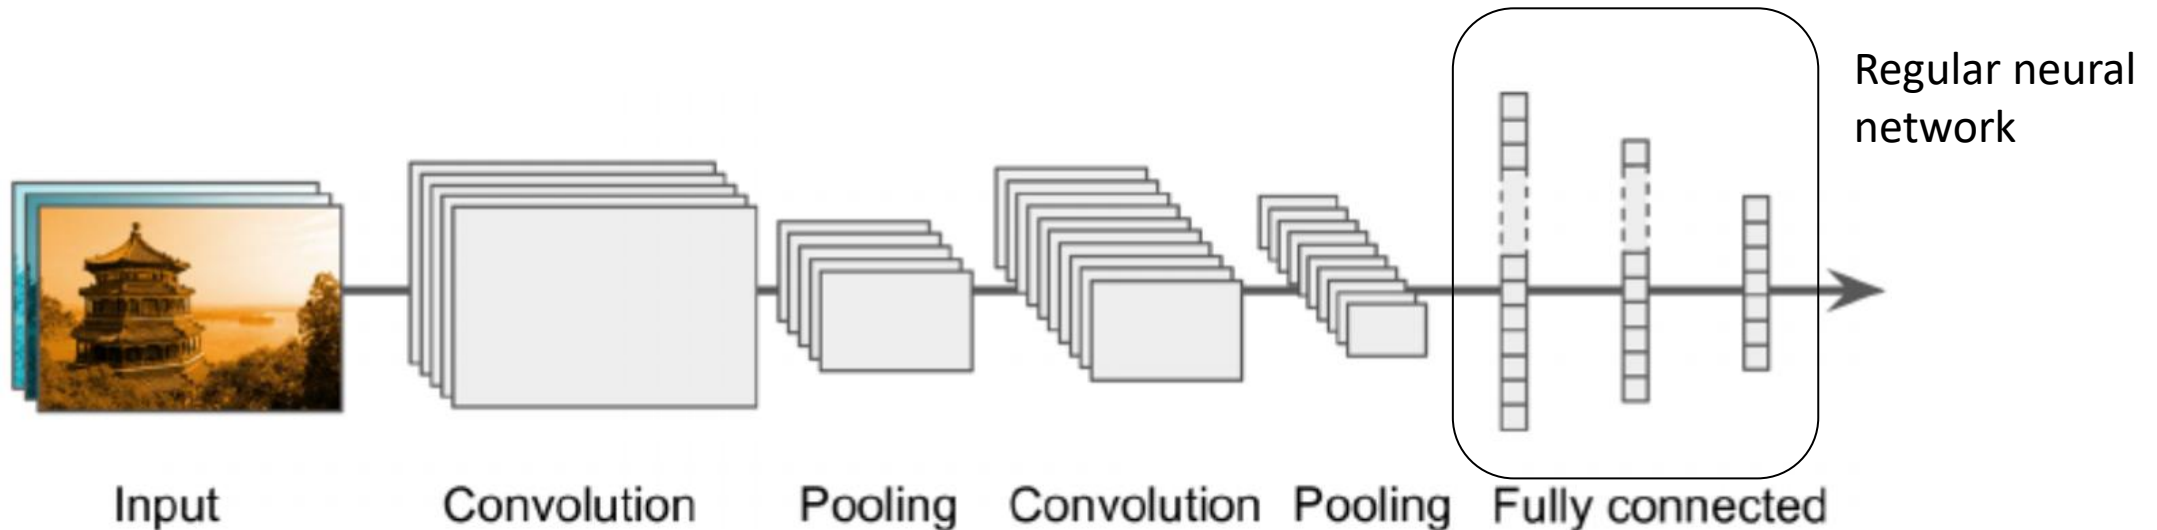
CNN Architectures

- Based on “stacking”
- Image gets smaller and smaller going through the layers
- Use a regular deep neural network at the end



CNN Architectures

- Can also stack two convolutional layers one after another
- Instead of having a 5 x 5 convo layer, have two 3 x 3 stacked convo layers. Might perform better!
- Double the filter size as you go deeper in CNN



LeNET-5

- Developed in 1998
- Used the MNIST data set (handwritten digits)



- **Labeled data**
- Black and white
- 28x28
- 60,000 images for training
- 10,000 images for testing

LeNET-5

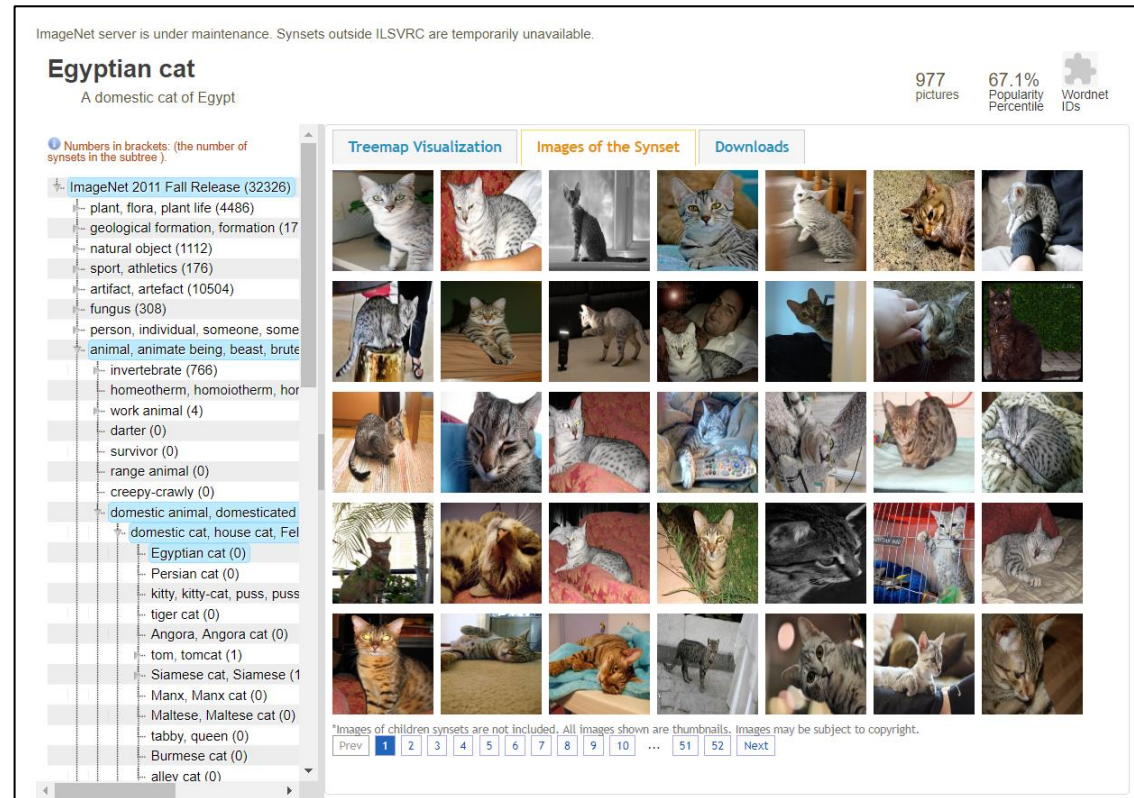
- Architecture:

Layer	Type	Maps	Size	Kernel size	Stride	Activation
Out	Fully connected	—	10	—	—	RBF
F6	Fully connected	—	84	—	—	tanh
C5	Convolution	120	1×1	5×5	1	tanh
S4	Avg pooling	16	5×5	2×2	2	tanh
C3	Convolution	16	10×10	5×5	1	tanh
S2	Avg pooling	6	14×14	2×2	2	tanh
C1	Convolution	6	28×28	5×5	1	tanh
In	Input	1	32×32	—	—	—

Error: less than 1%

AlexNET

- Developed in 2012
- Achieved 17% error rate in ILSVRC imagenet challenge
 - Labeled data (1,000 classes)
 - 150,000 images in total



AlexNET


- Architecture:

Layer	Type	Maps	Size	Kernel size	Stride	Padding	Activation
Out	Fully connected	—	1,000	—	—	—	Softmax
F10	Fully connected	—	4,096	—	—	—	ReLU
F9	Fully connected	—	4,096	—	—	—	ReLU
S8	Max pooling	256	6×6	3×3	2	valid	—
C7	Convolution	256	13×13	3×3	1	same	ReLU
C6	Convolution	384	13×13	3×3	1	same	ReLU
C5	Convolution	384	13×13	3×3	1	same	ReLU
S4	Max pooling	256	13×13	3×3	2	valid	—
C3	Convolution	256	27×27	5×5	1	same	ReLU
S2	Max pooling	96	27×27	3×3	2	valid	—
C1	Convolution	96	55×55	11×11	4	valid	ReLU
In	Input	3 (RGB)	227×227	—	—	—	—

- Used 50% dropout at F9 and 10
- Performed data augmentation
- Changed lighting conditions
- Local response normalization

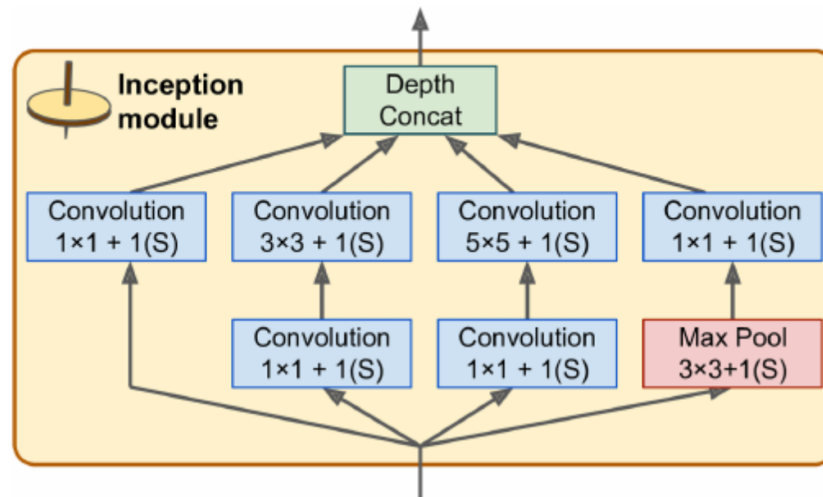
Data Augmentation

- Generate new images from existing training images
- Artificially increase the size
- Reduces overfitting



GoogLeNET

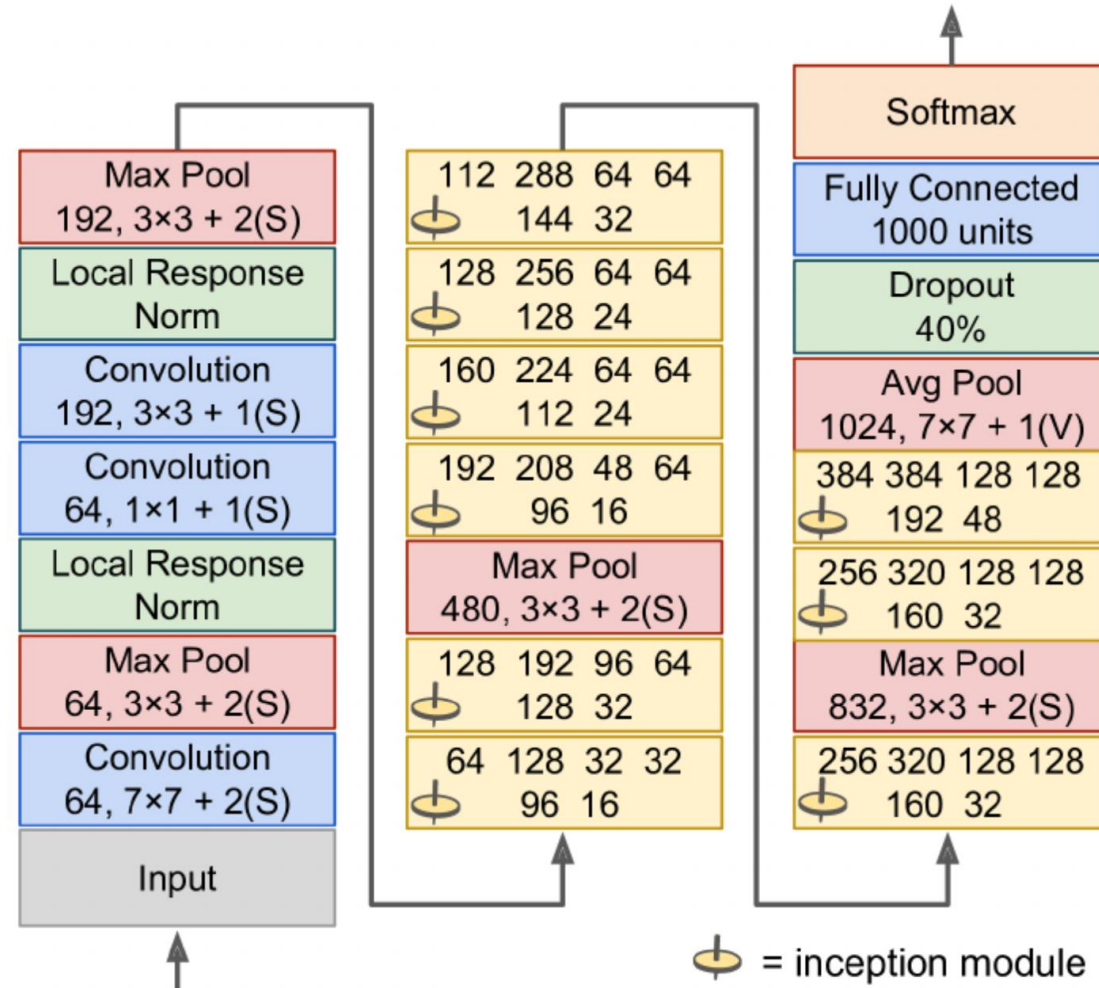
- Achieved an error rate of below 7% (ILSVRC imagenet challenge)
- Deeper network
- Reduced AlexNET's parameters by 10 times (from 60M to 6M)
- Used sub-networks called inception



Why use 1x1 filters?

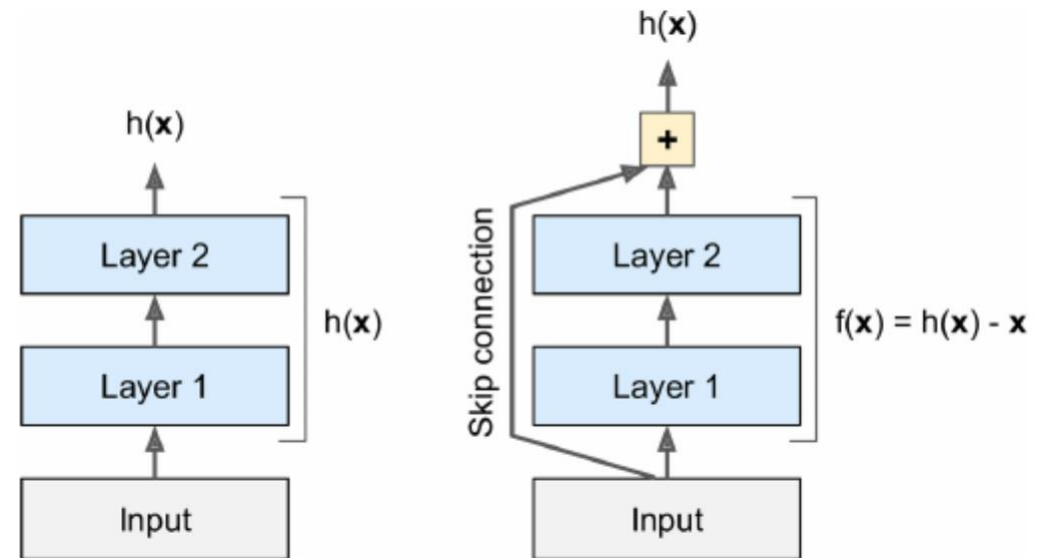
- Capture patterns along depth dimension
- Reduce dimensionality (bottleneck layers)

GoogLeNET

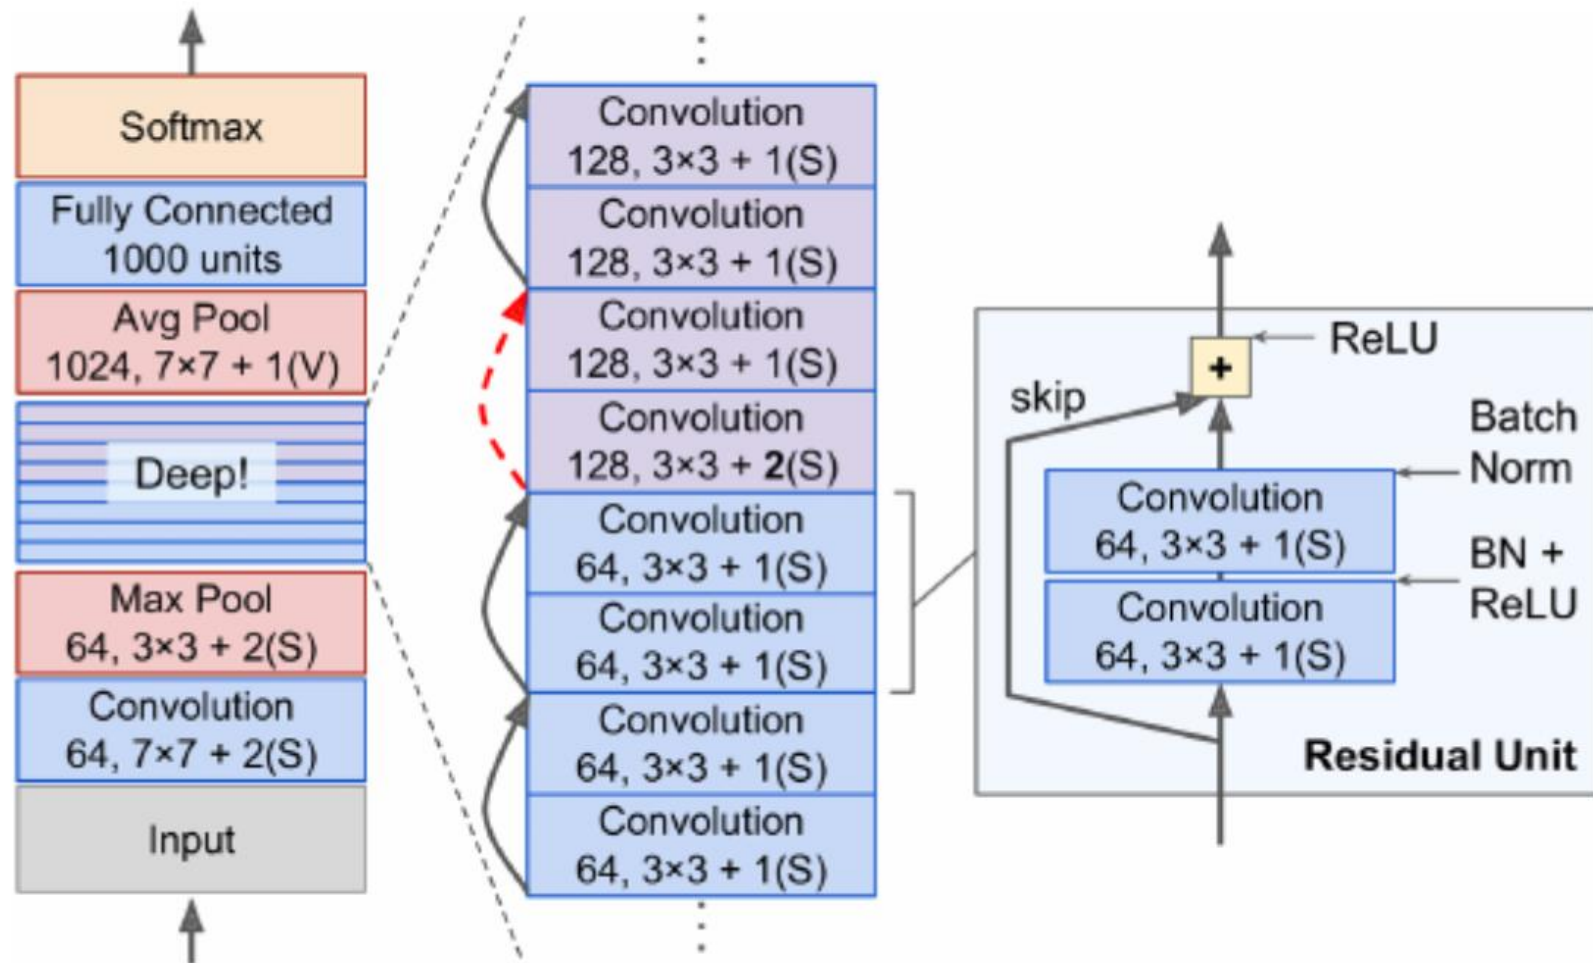


ResNET

- Reduced the error rate below 3.6%
- A total of 152 layers
- Used skip (shortcut) connections
 - Helps perform residual learning
- Skipping connections helps train faster



ResNET



Other Models

- VGGNet
- Xception (a variant of GoogLeNet)
- SENet (reduced the error rate to 2.25% in 2017)
- Inception v4
 - Combination of GoogLeNet and Resnet

Data Requirements

- We need **labeled** images (most of the time)





Image classification

New (unseen) image




OUTPUT


Cat

Object detection

New (unseen) image



OUTPUT



Cat 0.95

Object detection

- Example

A popular model:
You Only Look Once (YOLOv3)





Image segmentation

New (unseen) image



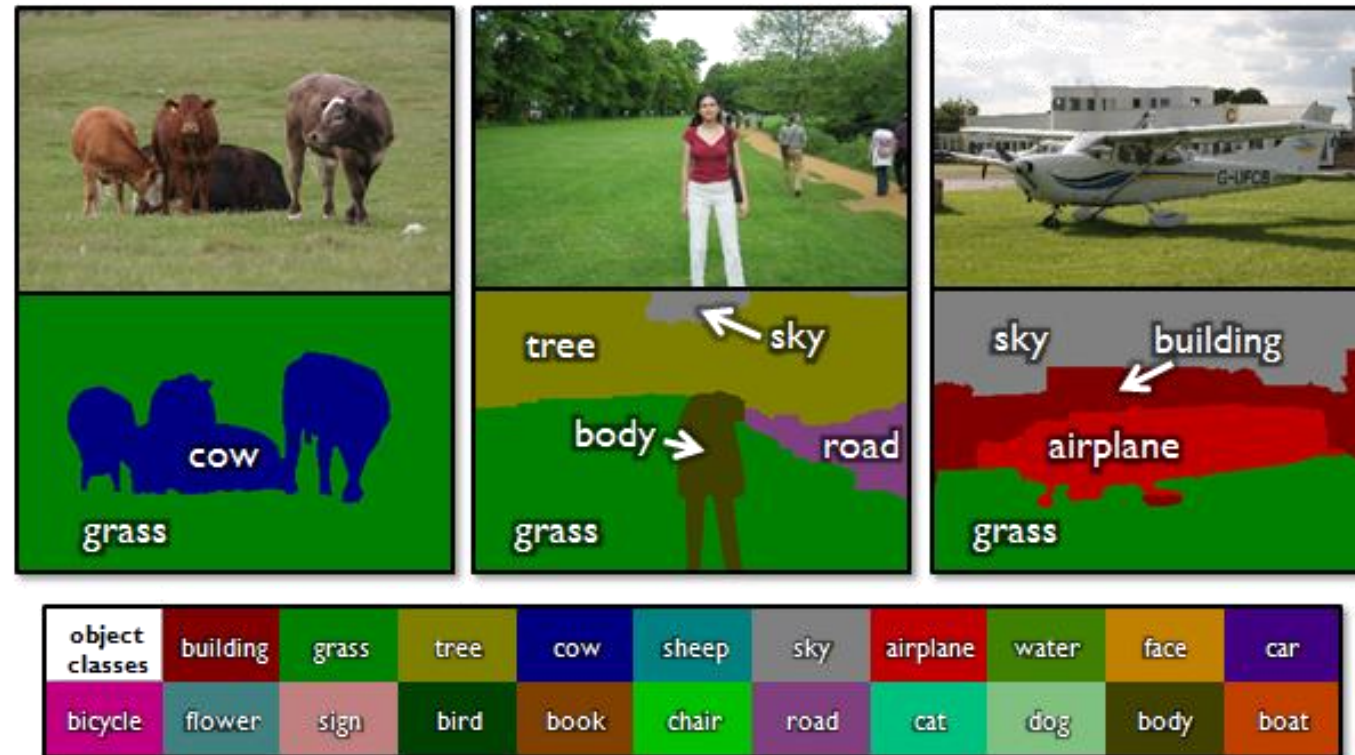
OUTPUT

Cat



What Kind of Output Does it Generate?

- Image segmentation example



How Are We Going to Use CNNs?

- Image classification

